Natural Language Processing Sequence Labeling and Hidden Markov Models

Felipe Bravo-Marquez

June 20, 2023

Overview

- The Sequence Labeling (or Tagging) Problem
- Generative models, and the noisy-channel model, for supervised learning
- Hidden Markov Model (HMM) taggers
    - Basic definitions
    - Parameter estimation
    - The Viterbi algorithm

This slides are based on the course material by Michael Collins: `http://www.cs.columbia.edu/˜mcollins/cs4705-spring2019/slides/tagging.pdf`

Sequence Labeling or Tagging Tasks

- Sequence Labeling or Tagging is a task in NLP different from document classification.
- Here the goal is to map a sentence represented as a sequence of tokens $x_1, x_2, \ldots, x_n$ into a sequence of tags or labels $y_1, y_2, \ldots, y_n$.
- Well known examples of this task are Part-of-Speech (POS) tagging and Named Entity Recognition (NER) to be presented next.

Part-of-Speech Tagging

INPUT: Profits soared at Boeing Co., easily topping forecasts on Wall Street, as their CEO Alan Mulally announced first quarter results.

OUTPUT: Profits/N soared/V at/P Boeing/N Co./N ,/, easily/ADV topping/V forecasts/N on/P Wall/N Street/N ,/, as/P their/POSS CEO/N Alan/N Mulally/N announced/V first/ADJ quarter/N results/N ./.

- N = Noun
- V = Verb
- P = Preposition
- Adv = Adverb
- Adj = Adjective
- ...

Named Entity Recognition

INPUT: Profits soared at Boeing Co., easily topping forecasts on Wall Street, as their CEO Alan Mulally announced first quarter results.

OUTPUT: Profits soared at [Company Boeing Co.], easily topping forecasts on [Location Wall Street], as their CEO [Person Alan Mulally] announced first quarter results.

Named Entity Extraction as Sequence Labeling

INPUT: Profits soared at Boeing Co., easily topping forecasts on Wall Street, as their CEO Alan Mulally announced first quarter results.

OUTPUT: Profits/NA soared/NA at/NA Boeing/SC Co./CC ,/NA easily/NA topping/NA forecasts/NA on/NA Wall/SL Street/CL ,/NA as/NA their/NA CEO/NA Alan/SP Mulally/CP announced/NA first/NA quarter/NA results/NA ./NA

- NA = No entity
- SC = Start Company
- CC = Continue Company
- SL = Start Location
- CL = Continue Location
- SP = Start Person
- CP = Continue Person

Our Goal

Training set:

1. Pierre/NNP Vinken/NNP ,/, 61/CD years/NNS old/JJ ,/, will/MD join/VB the/DT board/NN as/IN a/DT nonexecutive/JJ director/NN Nov./NNP 29/CD ./.

2. Mr./NNP Vinken/NNP is/VBZ chairman/NN of/IN Elsevier/NNP N.V./NNP ,/, the/DT Dutch/NNP publishing/VBG group/NN ./.

3. Rudolph/NNP Agnew/NNP ,/, 55/CD years/NNS old/JJ and/CC chairman/NN of/IN Consolidated/NNP Gold/NNP Fields/NNP PLC/NNP ,/, was/VBD named/VBN a/DT nonexecutive/JJ director/NN of/IN this/DT British/JJ industrial/JJ conglomerate/NN ./.

4. ...

Our Goal: From the training set, induce a function/algorithm that maps new sentences to their tag sequences.

Two Types of Constraints

Influential/JJ members/NNS of/IN the/DT House/NNP Ways/NNP and/CC Means/NNP Committee/NNP introduced/VBD legislation/NN that/WDT would/MD restrict/VB how/WRB the/DT new/JJ savings-and-loan/NN bailout/NN agency/NN can/MD raise/VB capital/NN ./.

"Local":

- e.g., "can" is more likely to be a modal verb MD rather than a noun NN

"Contextual":

- e.g., a noun is much more likely than a verb to follow a determiner

Sometimes these preferences are in conflict:

- The trash can is in the garage

Sequence Labeling as Supervised Learning

- We have a sequence of inputs $x = (x_1, x_2, \ldots, x_n)$ and corresponding labels $y = (y_1, y_2, \ldots, y_n)$.
- Task is to learn a function f that maps input sequences to label sequences: $f(x_1, x_2, \ldots, x_n) = y_1, y_2, \ldots, y_n$.
- We have a training set of labeled sequences: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \ldots, (x^{(m)}, y^{(m)})\}$.

Generative Approach for Sequence Labeling

- Generative models such as Naive Bayes can be used for sequence labeling tasks in NLP.
- Approach:
  - Training: Learn the joint distribution $p(x_1, x_2, \ldots, x_n, y_1, y_2, \ldots, y_n)$ of input sequences.
  - Decoding: Use the learned distribution to predict label sequences for new input sequences.
- Decoding in sequence labeling involves finding the label sequence with the highest joint probability: $\arg\max_{y_1, y_2, \ldots, y_n} p(x_1, x_2, \ldots, x_n, y_1, y_2, \ldots, y_n)$.

Hidden Markov Models

- HMMs provide a principled way to handle sequence labeling problems using generative modeling and efficient decoding algorithms.
- We have an input sentence $x = x_1, x_2, \ldots, x_n$ (x_i is the i-th word in the sentence).
- We have a tag sequence $y = y_1, y_2, \ldots, y_n$ (y_i is the i-th tag in the sentence).
- We'll use an HMM to define $p(x_1, x_2, \ldots, x_n, y_1, y_2, \ldots, y_n)$ for any sentence $x_1, \ldots, x_n$ and tag sequence $y_1, \ldots, y_n$ of the same length. [Kupiec, 1992]
- Then, the most likely tag sequence for x is:

$$\arg \max_{y_1, \ldots, y_n} p(x_1, \ldots, x_n, y_1, \ldots, y_n)$$

Trigram Hidden Markov Models (Trigram HMMs)

For any sentence $x_1, \ldots, x_n$ where $x_i \in V$ for $i = 1, \ldots, n$, and any tag sequence $y_1, \ldots, y_{n+1}$ where $y_i \in S$ for $i = 1, \ldots, n$, and $y_{n+1} = \text{STOP}$, the joint probability of the sentence and tag sequence is:

$$p(x_1, \ldots, x_n, y_1, \ldots, y_{n+1}) = \prod_{i=1}^{n+1} q(y_i | y_{i-2}, y_{i-1}) \prod_{i=1}^{n} e(x_i | y_i)$$

where we have assumed that $x_0 = x_{-1} = *$.

Parameters of the Model

- $q(s|u, v)$ for any $s \in S \cup \{\text{STOP}\}$, $u, v \in S \cup \{*\}$
- $e(x|s)$ for any $s \in S$, $x \in V$

An Example

If we have $n = 3$, x_1, x_2, x_3 equal to the sentence "the dog laughs", and y_1, y_2, y_3, y_4 equal to the tag sequence "D N V STOP", then:

$$p(x_1, \ldots, x_n, y_1, \ldots, y_{n+1}) = q(D|*, *) \times q(N|*, D)$$
$$\times q(V|D, N) \times q(STOP|N, V)$$
$$\times e(\text{the}|D) \times e(\text{dog}|N) \times e(\text{laughs}|V)$$

- STOP is a special tag that terminates the sequence.
- We take $y_0 = y_{-1} = *$, where $*$ is a special "padding" symbol.

Why the Name?

$$p(x_1, \ldots, x_n, y_1, \ldots, y_n) = q(\text{STOP}|y_{n-1}, y_n)$$
$$\times \prod_{j=1}^{n} q(y_j|y_{j-2}, y_{j-1})$$
$$\times \prod_{j=1}^{n} e(x_j|y_j)$$

• Markov Chain:

$$q(\text{STOP}|y_{n-1}, y_n) \times \prod_{j=1}^{n} q(y_j|y_{j-2}, y_{j-1})$$

• Observed:

$$e(x_j|y_j)$$

Smoothed Estimation

$$q(Vt|DT, JJ) = \lambda_1 \times \frac{\text{Count}(Dt, JJ, Vt)}{\text{Count}(Dt, JJ)}$$
$$+ \lambda_2 \times \frac{\text{Count}(JJ, Vt)}{\text{Count}(JJ)}$$
$$+ \lambda_3 \times \frac{\text{Count}(Vt)}{\text{Count}()}$$

where $\lambda_1 + \lambda_2 + \lambda_3 = 1$, and for all i, $\lambda_i \geq 0$.

$$e(\text{base}|Vt) = \frac{\text{Count}(Vt, \text{base})}{\text{Count}(Vt)}$$

Dealing with Low-Frequency Words

A common method is as follows:

- Step 1: Split vocabulary into two sets
  - Frequent words = words occurring $\geq$ 5 times in training
  - Low frequency words = all other words
- Step 2: Map low frequency words into a small, finite set, depending on prefixes, suffixes, etc.

Dealing with Low-Frequency Words: An Example

Below is an example of word classes for named entity recognition [Bikel et al., 1999]:

Word class	Example	Intuition
twoDigitNum	90	Two-digit year
fourDigitNum	1990	Four-digit year
containsDigitAndAlpha	$A8956 - 67$	Product code
containsDigitAndDash	$09 - 96$	Date
containsDigitAndSlash	$11/9/89$	Date
containsDigitAndComma	$23,000.00$	Monetary amount
containsDigitAndPeriod	1.00	Monetary amount, percentage
othernum	456789	Other number
allCaps	*BBN*	Organization
capPeriod	*M*.	Person name initial
firstWord	First word of sentence	No useful capitalization information
initCap	Sally	Capitalized word
lowercase	can	Uncapitalized word
other	,	Punctuation marks, all other words

Dealing with Low-Frequency Words: An Example

Original Sentence:
Profits/NA soared/NA at/NA Boeing/SC Co./CC ,/NA easily/NA topping/NA forecasts/NA on/NA Wall/SL Street/CL ,/NA as/NA their/NA CEO/NA Alan/SP Mulally/CP announced/NA first/NA quarter/NA results/NA ./NA

Transformed Sentence:
firstword/NA soared/NA at/NA initCap/SC Co./CC ,/NA easily/NA lowercase/NA forecasts/NA on/NA initCap/SL Street/CL ,/NA as/NA their/NA CEO/NA Alan/SP initCap/CP announced/NA first/NA quarter/NA results/NA ./NA

- NA = No entity
- SC = Start Company
- CC = Continue Company
- SL = Start Location
- CL = Continue Location
- SP = Start Person
- CP = Continue Person

Dynamic Programming

- Dynamic programming is a technique used to solve optimization problems by breaking them down into overlapping subproblems.
- It stores the solutions to these subproblems in a table, so they do not need to be recalculated.
- Dynamic programming can greatly improve the efficiency of algorithms.

Factorial

- Recursive implementation:

```python
def recur_factorial(n):
    # Base case
    if n == 1:
        return n
    else:
        return n * recur_factorial(n-1)
```

- Dynamic programming implementation:

```python
def dynamic_factorial(n):
    table = [0 for i in range(0, n+1)]

    # Base case
    table[0] = 1

    for i in range(1, len(table)):
        table[i] = i * table[i-1]
```

Fibonacci

- Recursive implementation:

```python
def recur_fibonacci(n):
    if n == 1 or n == 0:
        return 1
    else:
        return recur_fibonacci(n-1) + recur_fibo
```

- Dynamic programming implementation:

```python
def dynamic_fibonacci(n):
    table = [0 for i in range(0, n+1)]

    # Base case
    table[0] = 1
    table[1] = 1

    for i in range(2, len(table)):
        table[i] = table[i-1] + table[i-2]
```

Complexity

- Recursive factorial: Exponential complexity
- Dynamic factorial: Linear complexity
- Recursive Fibonacci: Exponential complexity
- Dynamic Fibonacci: Linear complexity

The Viterbi Algorithm

Problem: For an input $x_1 \ldots x_n$, find

$$\arg \max_{y_1 \ldots y_{n+1}} p(x_1 \ldots x_n, y_1 \ldots y_{n+1})$$

where the arg max is taken over all sequences $y_1 \ldots y_{n+1}$ such that $y_i \in S$ for $i = 1 \ldots n$, and $y_{n+1} = \text{STOP}$.
We assume that p takes the form:

$$p(x_1 \ldots x_n, y_1 \ldots y_{n+1}) = \prod_{i=1}^{n+1} q(y_i | y_{i-2}, y_{i-1}) \prod_{i=1}^{n} e(x_i | y_i)$$

Brute Force Search is Hopelessly Inefficient

Problem: For an input $x_1 \ldots x_n$, find

$$\arg \max_{y_1 \ldots y_{n+1}} p(x_1 \ldots x_n, y_1 \ldots y_{n+1})$$

where the arg max is taken over all sequences $y_1 \ldots y_{n+1}$ such that $y_i \in S$ for $i = 1 \ldots n$, and $y_{n+1} = \text{STOP}$.

The Viterbi Algorithm

The Viterbi algorithm efficiently computes the maximum probability of a tag sequence by using dynamic programming.

Steps:

- Define n as the length of the sentence.
- Define S_k for $k = -1 \ldots n$ as the set of possible tags at position k: $S_{-1} = S_0 = \{*\}$, $S_k = S$ for $k \in \{1 \ldots n\}$.
- Define
  $r(y_{-1}, y_0, y_1, \ldots, y_k) = \prod_{i=1}^{k} q(y_i | y_{i-2}, y_{i-1}) \prod_{i=1}^{k} e(x_i | y_i)$.
- Define a dynamic programming table: $\pi(k, u, v) =$ maximum probability of a tag sequence ending in tags u, v at position k.

An Example

$\pi(k, u, v)$ = maximum probability of a tag sequence ending in tags u, v at position k

The man saw the dog with the telescope

An Example

$S = \{D, N, P, V\}$

```
          The  man  saw  the  dog  with  the  telescope
*    *    1    2    3    4    5    6     7    8
```


D	D	D	D	D	P — D
N	N	N	N	N	
V	V	V	V	V	
P	P	P	P	P	

There are many possible
sequeces of tags.
Each of these will have a
probability calculated from
parameters q and e.

$\Pi(7, P, D) =$ This is the maxmimum probability of any
of those tag sequences ending in P D
at position 7, the path represents the
sequence with the maximum
probability.

A Recursive Definition

Base case:

$$\pi(0, *, *) = 1$$

Recursive definition: For any $k \in \{1 \ldots n\}$, for any $u \in S_{k-1}$ and $v \in S_k$:

$$\pi(k, u, v) = \max_{w \in S_{k-2}} \left(\pi(k-1, w, u) \times q(v|w, u) \times e(x_k|v) \right)$$
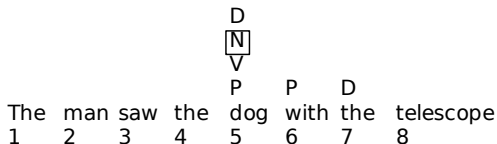
Justification for the Recursive Definition

For any $k \in \{1 \ldots n\}$, for any $u \in S_{k-1}$ and $v \in S_k$:

$$\pi(k, u, v) = \max_{w \in S_{k-2}} (\pi(k - 1, w, u) \times q(v|w, u) \times e(x_k|v))$$

The man saw the dog with the telescope

Justification for the Recursive Definition

```
                      D
                      N
                      V
                      P    P    D
  The  man  saw  the  dog  with the  telescope
  1    2    3    4    5    6    7    8
```

$$\mathcal{S}_5 = \mathcal{S} = \{D, N, V, P\}$$

$$\Pi(7, P, D) = \max_{w \in \mathcal{S}_5} \left(\Pi(6, w, P) \times q(D|w, P) \times e(\text{the}|D)\right)$$

If we think of any tag sequence that ends with tags P and D at position 7,
it must contain some tag at position 5.
We are basically searching for the tag that maximizes the probability at position 5.

The Viterbi Algorithm

Input: a sentence $x_1 \ldots x_n$, parameters $q(s|u, v)$ and $e(x|s)$.
Initialization: Set $\pi(0, *, *) = 1$.
Define $S_{-1} = S_0 = \{*\}$, $S_k = S$ for $k \in \{1 \ldots n\}$.
Algorithm:

- For $k = 1 \ldots n$,
- For $u \in S_{k-1}$, $v \in S_k$,

$$\pi(k, u, v) = \max_{w \in S_{k-2}} (\pi(k - 1, w, u) \times q(v|w, u) \times e(x_k|v))$$

- Return $\max_{u \in S_{n-1}, v \in S_n}(\pi(n, u, v) \times q(\text{STOP}|u, v))$

The Viterbi Algorithm

Algorithm 1: Viterbi Algorithm

Input: a sentence $x_1 \ldots x_n$, parameters $q(s|u,v)$ and $e(x|s)$

Initialization: Set $\pi(0, *, *) = 1$; $S_{-1} = S_0 = \{*\}$, $S_k = S$ for $k \in \{1 \ldots n\}$.

for $k = 1$ **to** n **do**

    for $u \in S_{k-1}, v \in S_k$ **do**

$$\pi(k, u, v) = \max_{w \in S_{k-2}} (\pi(k-1, w, u) \times q(v|w, u) \times e(x_k|v))$$

    end

end

`return` $(\max_{u \in S_{n-1}, v \in S_n}(\pi(n, u, v) \times q(STOP|u, v)))$

The Viterbi Algorithm with Backpointers

Algorithm 2: Viterbi Algorithm with Backpointers

Input: a sentence $x_1 \ldots x_n$, parameters $q(s|u, v)$ and $e(x|s)$
Initialization: Set $\pi(0, *, *) = 1$; $S_{-1} = S_0 = \{*\}$, $S_k = S$ for $k \in \{1 \ldots n\}$.

for $k = 1$ **to** n **do**
    for $u \in S_{k-1}, v \in S_k$ **do**

$$\pi(k, u, v) = \max_{w \in S_{k-2}} \left(\pi(k - 1, w, u) \times q(v|w, u) \times e(x_k|v) \right)$$

$$\text{bp}(k, u, v) = \arg \max_{w \in S_{k-2}} \left(\pi(k - 1, w, u) \times q(v|w, u) \times e(x_k|v) \right)$$

    end
end

$(y_{n-1}, y_n) = \arg \max_{(u,v)} (\pi(n, u, v) \times q(\text{STOP}|u, v))$;   // Find maximum
 probability and corresponding tags
for $k = (n - 2)$ **to** 1 **do**
    $y_k = \text{bp}(k + 2, y_{k+1}, y_{k+2})$;   // Retrieve tag sequence using
    backpointers
end

return(*the tag sequence* $y_1 \ldots y_n$) ; // Return the final tag sequence

---

The Viterbi Algorithm: Running Time

- $O(n|S|^3)$ time to calculate $q(s|u,v) \times e(x_k|s)$ for all k, s, u, v.
- $n|S|^2$ entries in π to be filled in.
- $O(|S|)$ time to fill in one entry.

$\Rightarrow O(n|S|^3)$ time in total.

Pros and Cons

- Hidden Markov Model (HMM) taggers are simple to train (compile counts from training corpus).
- They perform relatively well (over 90
- Main difficulty is modeling e(word|tag), which can be very complex if "words" are complex.

Questions?

Thanks for your Attention!

References I

Bikel, D. M., Schwartz, R. M., and Weischedel, R. M. (1999).
An algorithm that learns what's in a name.
Mach. Learn., 34(1-3):211–231.

Kupiec, J. (1992).
Robust part-of-speech tagging using a hidden markov model.
Computer speech & language, 6(3):225–242.